

Lean Introduction

From formal proofs to basic tactics

Today's route

Kernel checks proofs



Types organize data



Tactics build proof terms

July 15, Tuesday

Type judgments,
propositions-as-types, basic
tactics, and AI-assisted Lean
work.

Learning Goals

Read Lean expressions

Use `a : A` and `#check`; separate type checking from evaluation.

Understand proofs as terms

Curry–Howard: propositions are types; proofs are terms.

Write small programs

Define functions with `def`; inspect them with `#eval`, `#reduce`, and `#print`.

Build proof terms with tactics

Use `intro`, `exact`, `apply`, `constructor`, `cases`.


Use AI carefully

Let AI propose candidates; verify names, statements, and scripts in Lean.

Reason by equality

Use `rfl`, `rw`, and `simp` for rewriting.

Four-Part Outline



First: why a kernel can check proofs and why AI output still needs checking.

Then: the type-theoretic language in which Lean states and checks goals.

Six-Day Course Overview

Day	Date	Topic
D1	Jul 15 (Tue)	Types, Prop, basic tactics, AI tools
D2	Jul 16 (Wed)	Inductive types, structures, type classes
D3	Jul 17 (Thu)	Instances, AI workflows, algebra, SimpleGraph
D4	Jul 18 (Fri)	Analysis, topology, filters, limits
D5	Jul 19 (Sat)	Functional programming, Monad
D6	Jul 20 (Sun)	Algorithm verification, termination, complexity

Tutorial and Project Schedule

Afternoon tutorials

- ▶ D1: environment setup; AI tool configuration; Mathlib search; `ring/linarith/decide`
- ▶ D2: inductive, structure, class exercises
- ▶ D3: algebra/combinatorics exercises; project selection and grouping
- ▶ D4–D5: project work; AI drafting, Mathlib search, Lean repair
- ▶ D6: common feedback on spec, invariants, termination, complexity


Project timeline

- ▶ D3: form groups, select topics
- ▶ Weekend D4–D5: independent group work with AI tools
- ▶ D6: whole-class feedback and wrap-up


Part I: Formal Proofs and AI

Three questions


1. Why does Lean know whether a proof is correct?
2. Where does Lean fit in the history of mathematics?
3. How should we use AI tools alongside Lean?



Search, Proof, and Verification Working Together




Donald Knuth



In 2023–2025, a Knuth problem was resolved by combining search, explanation, computation, and proof checking.

From Recording to Checkable Proofs



Lean's position

A proof written in Lean is an object. The kernel checks that object step by step against explicit rules.

Ancient Mathematics: Recording and Computing

Record

quantities, land, debt,
astronomy

Compute

apply a rule to obtain a
result


Explain

ask why the result must
hold

Early mathematical traditions served memory and calculation first. Formalization inherits the same discipline:

- ▶ records must be accurate;
- ▶ computations must be reliable;
- ▶ reasons must be inspectable.

Euclid: Proofs as Dependency Networks



Elements (c. 300 BCE) established a model still in use:

- ▶ start from definitions and postulates;
- ▶ order propositions by dependency;
- ▶ use diagrams for intuition while proofs carry the conclusion.

Lean's dependency graph is a direct descendant.

Al-Khwarizmi: Algorithms and Knowledge Preservation

Algebra

Transform problems into symbols and rules.

Algorithm

Give steps that can be executed repeatedly.

Mathlib inherits the tradition of organizing steps and knowledge for reuse.

Modern Proofs Grow Long: Fermat to Wiles


Fermat's Last Theorem:

$a^n + b^n = c^n$, $n > 2$ has no positive integer solution.

- ▶ Short to state; centuries to prove.
- ▶ Wiles (1995): semi-stable case via modular forms and Galois representations.
- ▶ Complete modularity theorem completed by subsequent work.

As proofs grow longer,
the question of *who*
checked each step
becomes pressing.

When Verification Becomes Part of the Story



The community needs not only answers, but a form of expression that can be jointly inspected.

Hilbert: Can Proof Checking Be Mechanized?

David Hilbert (1862–1943) aimed to:

- ▶ place mathematics on an explicit formal foundation;
- ▶ express proofs in a clear symbolic language with fixed rules;
- ▶ prove the consistency of those foundations.

Machine-checked proofs are a natural endpoint of this program: instead of trusting authority or intuition, every step is held to explicit rules.

Principia Mathematica: Reasoning as Symbols

Russell and Whitehead (1910–1913) aimed to derive mathematics from explicit symbolic rules.

- ▶ Famous: $1 + 1 = 2$ required hundreds of pages.
- ▶ Demonstrated the power of formalization — and its writing cost.
- ▶ The cost motivated the search for automation.

Gödel: Limits Are Boundaries, Not Endpoints

Gödel's incompleteness theorems (1931):

- ▶ A sufficiently strong consistent formal system cannot prove all true statements.
- ▶ Such a system typically cannot prove its own consistency internally.

This does not prevent us from mechanically checking concrete proofs within a fixed system.

Lean is not a machine that proves all truths. It is a tool that checks a given formal proof.

Two Routes: Proof Search vs. Proof Checking

Proof search

Machine tries to find a derivation.

- ▶ 1956: Logic Theorist
- ▶ Today: ATP solvers, AI provers

Proof checking

The proof object is supplied; the kernel verifies it.

- ▶ 1968: Automath
- ▶ Today: Coq, Lean kernel

In Lean, the two meet: tactics and AI help with search, but every proof must pass the kernel.

Automath: Proofs as Formal Objects

Nicolaas de Bruijn's Automath (1968):

- ▶ First system to express proofs as precisely typed objects.
- ▶ Every line is checked against a rule.
- ▶ A small trusted kernel suffices.

The key shift: from “reading a text” to “checking an object.” This is exactly what the Lean kernel does.

Four-Color Theorem: Computer-Assisted Proof

- ▶ 1976: Appel and Haken prove the four-color theorem using computer case analysis.
- ▶ Question: can the mathematical community trust computer output?
- ▶ 2005: Georges Gonthier's team completes a machine-checked proof in Coq.


This binds two questions together: “Is the computation correct?” and “Is the proof checkable?”

Flyspeck: Engineering a Complex Proof

- ▶ 1611: Kepler's sphere-packing conjecture.
- ▶ 1998: Thomas Hales announces a computer-assisted proof.
- ▶ 2003–2014: Flyspeck project formalizes it in HOL Light and Isabelle.


A modern proof is often a project: the target theorem is at the top; beneath it is a large body of reusable infrastructure. Formalization is not a one-time translation; it is long-term library construction.

Proof Assistants as Infrastructure



Proof assistants moved from individual checkers to shared mathematical infrastructure.

The Trust Chain of Formal Proofs



Lean checks whether the formalized statement has a valid proof. Mathematicians separately verify that the statement captures the intended mathematics.

Lean: Language, Checker, and Math Library

Language

write programs and
definitions

Checker

kernel checks proof terms

Library

Mathlib stores reusable
mathematics

The working mindset: enter a reusable, checkable mathematical knowledge base rather than proving the world from scratch.

Mathlib: A Reusable Math Library

Mathlib contains:

- ▶ **Definitions** — types, structures, functions;
- ▶ **Theorems** — results with machine-checked proofs;
- ▶ **Notation** — mathematical symbols;
- ▶ **Type classes** — algebraic, order, topological structures.

Library-centered workflow: find existing definitions and theorems first; then write your own proof on top.

Theorem Names Are Searchable Addresses

Lean theorem names follow a pattern:

- ▶ **prefix** — Nat, Int, Set (namespace / domain);
- ▶ **core word** — add, mem, ext (operation or relation);
- ▶ **modifier** — comm, left (direction or property).

```
#check Nat.add_comm      -- Nat.add_comm : a + b = b + a
#check Int.add_comm
#check Set.mem_union
#check Finset.mem_union
```

Live Lean: [codeD1/D1.lean](#)

Different Namespaces, Same Short Name

Many mathematical objects have an “extensionality” theorem named `ext`; namespaces keep them distinct.

```
#check Set.ext
-- (h : forall x, x in s <-> x in t) : s = t

#check Finset.ext
-- (h : forall x, x in A <-> x in B) : A = B
```

This is why namespaces matter: short names can repeat; full names remain unique.

Live Lean: [codeD1/D1.lean](#)


File Organization Commands

```
import Mathlib          -- load a file or library
namespace Demo
  def double (n : Nat) : Nat := n + n
end Demo
section
  variable (n : Nat)
  example : Demo.double n = n + n := by rfl
end
open Nat
#check zero_add          -- same as Nat.zero_add
```

`import`: load a library; `namespace`: add a name prefix; `section`: scope local variables; `open`: drop a prefix temporarily.

Live Lean: [codeD1/D1.lean](#)

AI for Math: Where Does Lean Fit?



Where this course fits

We learn the target language: the precise Lean statements and proof scripts that AI tools must eventually produce and that formalization projects share.

Choosing AI for Lean Work

Proof routes

Ask for lemmas and a proof outline; do not trust a complete script blindly.

Code repair

Provide the error, goal state, imports, and a narrow edit target.

Retrieval

Use search and citations; verify names, versions, and statements in Lean.

Local deploy

Use open-weight models for privacy or batch experiments.

Operational rule: AI proposes candidates; students decompose; Lean gives the final verdict.

What AI Is Good at in Math Learning

Explore

examples, conjectures,
proof routes

Explain

definitions, analogies, small
examples

Collaborate

drafts, code completion,
gap finding

Use it as a co-pilot

AI broadens the search space and speeds up drafts. It does not replace the obligation to understand the statement and check the proof.

Hallucinations: The Dangerous Failure Mode

Looks plausible

Fluent proof text can contain fabricated theorem names, missing lemmas, or changed hypotheses.

Root cause

Language models optimize for plausible continuation, not for a global dependency graph.

Failure signs

Names fail under `#check`; goals silently change; key cases are skipped.

Response

Treat AI output as a draft candidate, not as mathematical evidence.


Making AI Output Checkable



Precise statements expose variables, assumptions, and quantifiers.

`#check`, `exact?`, `Loogle`, and `Lean errors` turn guesses into verifiable work.

Part II: Types and Functions



Read Lean from right to left when needed: first identify the type, then understand the expression inhabiting it.

Start Writing Lean

```
import Mathlib

#check Nat      -- Nat : Type
#eval 2 + 3     -- 5

example (P : Prop) : P -> P := by
  intro h
  exact h
```

Core habit: after every line, read the Infoview panel to see the current type or goal.

Live Lean: `codeD1/D1.lean`

Three Core Ideas of Type Theory

Judgment

$2 : \text{Nat}$

expression belongs to a type

Arrow

$f : A \rightarrow B$

input type to output type

Proof

$h : P$

evidence inhabits a
proposition

Main shift

Lean proofs are not natural-language paragraphs. They are objects built in type theory and checked by the kernel.

Type Judgments: $a : A$

The basic sentence in Lean is a *type judgment*: $a : A$, read “a has type A.”

```
#check 42           -- 42 : Nat
#check Nat          -- Nat : Type
#check Type         -- Type : Type 1
#check 42 + 1       -- 42 + 1 : Nat
-- #check 42 + "hi" -- error: type mismatch
```

A type error does not mean False; it means the expression was rejected before truth or proof was even considered.

Live Lean: [codeD1/D1.lean](#)

Every Expression Has a Type

```
#check Nat      -- Nat : Type
#check 37       -- 37 : Nat
#check "hello"  -- "hello" : String
#check true     -- true : Bool
```

- ▶ #check asks for the type.
- ▶ #eval asks for the computed value.
- ▶ These are two different questions.

Live Lean: `codeD1/D1.lean`

#check, #reduce, #print

```
def onePlusOne : Nat := 1 + 1

#check onePlusOne    -- onePlusOne : Nat
#reduce onePlusOne    -- 2
#print onePlusOne     -- def onePlusOne : Nat := 1 + 1
```

#check What type does this name or expression have?

#reduce Reduce the expression to its normal form.

#print What is stored in the environment for this name?

Live Lean: [codeD1/D1.lean](#)

Types Also Have Types: Universes

```
#check Nat      -- Nat : Type
#check Prop     -- Prop : Type
#check Type     -- Type : Type 1
#check Type 1   -- Type 1 : Type 2
-- Sort 0 = Prop, Sort 1 = Type, Sort u : Type (u+1)
```

Why layering? If there were one “type of all types” $U : U$, self-reference would reintroduce paradoxes. Lean uses $\text{Type } u : \text{Type } (u+1)$ to avoid them.

Live Lean: [codeD1/D1.lean](#)

Why Not $U : U$? Russell's Problem

The naive idea: a “set of all sets that do not contain themselves.”

$$R = \{x \mid x \notin x\} \quad \text{Is } R \in R?$$

Both answers lead to contradiction. Lean avoids this with universe levels:

```
#check Type      -- Type : Type 1
#check Type 1    -- Type 1 : Type 2
-- Type u : Type (u + 1)
```

Live Lean: [codeD1/D1.lean](#)

Functions: Basic Building Blocks

```
def double (n : Nat) : Nat :=  
  n + n  
  
#check double    -- double : Nat -> Nat  
#eval double 21 -- 42
```

Parts of a definition:

- ▶ `def` — give the object a name;
- ▶ `(n : Nat)` — an explicit input parameter;
- ▶ `: Nat` — the return type;
- ▶ `n + n` — the function body.

Live Lean: `codeD1/D1.lean`

Anonymous Functions and Higher-Order Functions

```
#check fun x : Nat => x + 1

def applyTwice (f : Nat -> Nat) (x : Nat) : Nat :=
  f (f x)

#eval applyTwice (fun n => n + 1) 10  -- 12
```

A function can take another function as input. This is the core of functional programming.

Live Lean: [codeD1/D1.lean](#)

Multi-Argument Functions: Currying

```
def add3 (a b c : Nat) : Nat := a + b + c

#check add3          -- add3 : Nat -> Nat -> Nat -> Nat
#check add3 1        -- add3 1 : Nat -> Nat -> Nat
#check add3 1 2      -- add3 1 2 : Nat -> Nat
#check add3 1 2 3    -- add3 1 2 3 : Nat
```

`add3 1` is still a function waiting for two more arguments. Partial application comes for free in Lean.

Live Lean: [codeD1/D1.lean](#)

Implicit Arguments: Lean Infers Them

```
def myConst {A B : Type} (a : A) (_ : B) : A := a

#eval myConst 5 "ignored" -- 5
```

Curly-bracket parameters `{A B : Type}` are inferred by Lean from the surrounding context.

When starting out, read implicit arguments as “type information that Lean fills in automatically.”

Live Lean: [codeD1/D1.lean](#)

Three Bracket Types

The bracket style on a parameter tells Lean *who provides it*.

`(x :  $\alpha$ )` **Explicit**. You supply it. Used for the main mathematical objects.

`{ $\alpha$  : Type}` **Implicit**. Lean infers it from later arguments or the goal.

`[Monoid  $\alpha$ ]` **Instance**. Lean searches the instance database (algebraic structure, order, topology, ...).

```
#check List.length      -- hides inferred arguments
#check @List.length     -- shows all arguments
#check @Nat.add_comm    -- shows [AddCommMonoid] instance
```

Live Lean: [codeD1/D1.lean](#)

Calling Functions with Implicit Arguments

```
def same {alpha : Type} (x : alpha) : alpha := x

#check same 3          -- same 3 : Nat  (alpha inferred)
#check @same Nat 3      -- explicit: alpha = Nat


example {alpha : Type} [AddMonoid alpha] (x : alpha) :
  x + 0 = x := by simp

example : same (alpha := Nat) 3 = 3 := by rfl
```

- ▶ @f shows all hidden arguments.
- ▶ f (alpha := Nat) supplies a named implicit argument.
- ▶ [AddMonoid alpha] triggers instance search.

Live Lean: [codeD1/D1.lean](#)

Part III: Curry-Howard and Prop



What tactics do

A tactic script is a convenient way to construct a term whose type is the current proposition.

Dependent Functions: Π , $\forall$, and fun

```
#check ((n : Nat) -> Fin (n + 1))
-- the output type depends on the input n

#check (forall n : Nat, n = n)
-- universally quantified proposition

example : (forall n : Nat, n = n) := by
  intro n
  rfl
```

Read $\forall n, P\ n$ as “a function type: give any n , return a proof of $P\ n$.” Then `intro` is simply the function-argument step.

Live Lean: [codeD1/D1.lean](#)

Propositions as Types

Data side

`42 : Nat`

`true : Bool`

a value inhabits a data type

Proof side

`h : 2 + 2 = 4`

`h : P`

a proof inhabits a proposition

The same colon does both jobs: “has type” and “is a proof of”.

The Curry-Howard Correspondence

Logic	Type theory
proposition P	type P
proof of P	term $p : P$
proof of $P \rightarrow Q$	function $f : P \rightarrow Q$
use hypothesis $h : P$	use term h

Mental model

To prove an implication, build a function. To use an implication, apply that function to a proof of its premise.

Per Martin-Löf (1942–)



Per Martin-Löf

Type theory as foundations

Martin-Löf Type Theory (1972–1984) is a direct ancestor of Lean's type system.

Propositions as types

This is not just an analogy: a proposition is treated as the type of its proofs.

Dependent types

Types may depend on values, e.g. vectors indexed by their length.

Curry-Howard Dictionary

Logic	Proof construction	What to do
True	<code>True.intro</code>	trivial; one proof
False	no constructor	nothing proves it
$P \rightarrow Q$	<code>fun hP => hQ</code>	give P, return Q
$P \wedge Q$	<code>And.intro hP hQ</code>	supply both proofs
$P \vee Q$	<code>Or.inl hP</code> or <code>Or.inr hQ</code>	supply one side
not P	<code>fun h => absurd</code>	$P \rightarrow \text{False}$

Curry-Howard in Lean

```
example (P Q : Prop) : P -> Q -> And P Q := by
  intro hP
  intro hQ
  exact And.intro hP hQ


#check And.intro
-- And.intro : P -> Q -> P /\ Q
```

- ▶ `intro hP` puts the proof of `P` into context.
- ▶ `And.intro hP hQ` assembles both proofs into a proof of `And P Q`.

Tactics are not magic commands; they construct a term whose type is the current goal.

Live Lean: [codeD1/D1.lean](#)

Part IV: Logic Tactics



What tactics are

Commands that construct a proof term step by step.

How to debug

Move cursor line by line and verify each goal transformation.

Reading the Lean Infoview

The Infoview has two regions: available facts above the line, target below.

```
P Q : Prop          -- context (hypotheses)
hP : P
hPQ : P -> Q
----- Goal -----
Q                   -- what you need to prove
```

Context

names and types you may use

Goal

the proposition still to prove

exact: Closing a Goal

```
example (P : Prop) (h : P) : P := by
  exact h
```

Infoview reading:

- ▶ Context: $h : P$
- ▶ Goal: P
- ▶ `exact h` closes the goal because h has exactly type P .

`exact` means: I already have a term whose type matches the goal.

Live Lean: [codeD1/D1.lean](#)

intro: Proving Implications

```
example (P : Prop) : P -> P := by
  intro h
  exact h
```

Goal changes:

1. Before `intro h`: goal is $P \rightarrow P$
2. After `intro h`: context gains $h : P$; goal is P
3. `exact h` closes it.

`intro` moves an assumption from the goal into the context.

Live Lean: [codeD1/D1.lean](#)

apply: Using a Theorem Backwards

```
example (P Q : Prop) (hPQ : P -> Q) (hP : P) : Q := by
  apply hPQ
  exact hP
```

- ▶ Goal before apply hPQ: Q
- ▶ Goal after: P
- ▶ exact hP closes it.

apply reduces the current goal to the premises of a known implication.

Live Lean: [codeD1/D1.lean](#)

constructor: Splitting a Conjunction

```
example (P Q : Prop) (hP : P) (hQ : Q) : And P Q := by
  constructor
  exact hP
  exact hQ
```

constructor on `And P Q` creates two subgoals:

1. After constructor: goal 1 is `P`, goal 2 is `Q`.
2. `exact hP` closes goal 1.
3. `exact hQ` closes goal 2.

Live Lean: [codeD1/D1.lean](#)

cases: Splitting a Disjunction

```
example (P Q : Prop) : Or P Q -> Or Q P := by
  intro h
  cases h with
  | inl hP => exact Or.inr hP
  | inr hQ => exact Or.inl hQ
```

cases h on Or P Q creates two branches:

- ▶ | inl hP => branch: context has hP : P
- ▶ | inr hQ => branch: context has hQ : Q

Each branch must close the goal independently.

Live Lean: [codeD1/D1.lean](#)

Negation, False, and contradiction

```
example (P : Prop) (hP : P) (hNotP : Not P) : False := by
  contradiction
```

```
example (P : Prop) (hP : P) (hNotP : P -> False) : False := by
  contradiction
```

- ▶ Not P is defined as $P \rightarrow \text{False}$.
- ▶ False has no constructor; nothing proves it.
- ▶ When context contains both P and Not P, contradiction finds the conflict automatically.

Live Lean: [codeD1/D1.lean](#)

use: Witness for Existence

```
example (m : Nat) : Exists (fun n : Nat => 0 + n = m) := by
  use m
  simp
```

- ▶ use m provides the concrete witness m.
- ▶ The remaining goal is the property: $0 + m = m$.
- ▶ simp closes it.

Exists is written $\exists$ in Lean: $\exists n, 0 + n = m$.

A proof of an existential = a concrete witness + a proof that it satisfies the property.

Live Lean: `codeD1/D1.lean`

Definitional Equality: What rfl Checks

```
def onePlusOne : Nat := 1 + 1

#reduce (fun n : Nat => n + 1) 3  -- 4
#reduce 2 + 3                    -- 5
#reduce onePlusOne               -- 2
```

Two expressions are *definitionally equal* if the Lean kernel reduces both to the same normal form. `rfl` closes a goal exactly when both sides reduce the same way.

Live Lean: [codeD1/D1.lean](#)

rfl: Same After Unfolding

```
example (a : Nat) : a + 0 = a := by
  rfl                                -- works: definitionally equal

example (a : Nat) : 0 + a = a := by
  exact Nat.zero_add a             -- needs a theorem
```

- ▶ $a + 0 = a$: Lean reduces both sides to the same thing.
- ▶ $0 + a = a$: reduction does not equate them; the theorem `Nat.zero_add` is needed.

Live Lean: [codeD1/D1.lean](#)

When rfl Fails

```
example : 2 + 3 = 5 := by rfl    -- succeeds

-- example : 2 + 2 = 5 := by rfl -- fails: not equal

example (a : Nat) : 0 + a = a := by
  exact Nat.zero_add a
```

When rfl fails:

- ▶ Use #reduce to see what each side computes to.
- ▶ Search for a theorem using the goal's shape.
- ▶ Use rw or simp to reduce it to something rfl can handle.

Read rfl as “please compare by definition,” not “search for a proof.”

Live Lean: [codeD1/D1.lean](#)

rw: Rewriting with an Equation

```
example (a b c : Nat) (h : a = b) : a + c = b + c := by
  rw [h]          -- replaces a with b in the goal
```

```
example (a b c : Nat) (h : a = b) : c + b = c + a := by
  rw [h.symm]     -- replaces b with a in the goal
```

- ▶ `rw [h]` rewrites left-to-right using `h`.
- ▶ `rw [h.symm]` or equivalently `rw [ h]` rewrites right-to-left (is the idiomatic Lean 4 form).
- ▶ `rw` gives precise control over exactly which equality is used.

Live Lean: `codeD1/D1.lean`

simp: Automatic Simplification

```
example (a : Nat) : a + 0 = a := by
  simp

def square (n : Nat) : Nat := n * n

example : square 3 = 9 := by
  simp [square]
```

- ▶ `simp` applies a library of standard simplification lemmas.
- ▶ `simp [f]` additionally unfolds the definition `f`.
- ▶ `simp` is automatic and fast; it may close goals that would need several `rw` steps.

Live Lean: `codeD1/D1.lean`

rw vs simp: When to Use Which

	rw [h]	simp / simp [h]
When	you know exactly which equation to use	standard cleanup or goal contains simp lemmas
Control	precise, one step	automatic
Speed	slower to write	fast
Use for	targeted rewrites	finishing off a goal

Two More Useful Tactics: omega and decide

```
-- omega: linear arithmetic over Nat and Int
example (n : Nat) (h : n > 3) : n >= 4 := by omega

example (a b : Nat) : a + b = b + a := by omega

-- decide: any decidable proposition with finite cases
example : 2 + 3 = 5 := by decide

example : (3 : Nat) < 7 := by decide
```

- ▶ omega solves linear arithmetic goals automatically — addition, subtraction, inequalities over Nat/Int.
- ▶ decide evaluates a decidable proposition at compile time; works whenever the kernel can reduce the goal to True.
- ▶ Use these before reaching for rw or simp on numeric goals.

Live Demo Path

The session-opening code snippet follows this route:

```
import Mathlib
#check Nat          -- step 1: see a type
#eval 2 + 3         -- step 1: compute a value

def double (n : Nat) : Nat := n + n -- step 2: write a function
#eval double 21

example (P : Prop) : P -> P := by   -- step 3: enter proof mode
  intro h
  exact h
```

First check types, then define functions, then prove.

Live Lean: `codeD1/D1.lean`

Nine Actions to Remember

Inspect

#check e
#eval e

Transform

intro h
apply h
rw [h]

Automate

simp
omega
decide

Close

exact h
rfl

Split

constructor
cases h
use t

Habit

read goal
choose action
confirm change

Suggested Exercises

Work through these in `lean-code/D1.lean`. After each tactic line, record how the Infoview goal changes.

```
-- 1. implication chaining (intro, apply, exact)
example (P Q R : Prop) : (P -> Q) -> (Q -> R) -> P -> R := by
  sorry

-- 2. rewriting (rw, omega)
example (a b : Nat) (h : a = b) : a + a = b + b := by
  sorry

-- 3. And commutativity (intro, cases, constructor, exact)
example (P Q : Prop) : P /\ Q -> Q /\ P := by
  sorry

-- 4. existential witness (use, omega)
example : Exists (fun n : Nat => n > 10) := by
  sorry
```

Learning Lean is not memorizing commands; it is learning to *read the goal, choose an action, and confirm the goal changes*.