

D2: Inductive Types and Classes

Constructors, recursion, induction, structures, and type classes

Summer School on Lean

July 14, Tuesday

Learning Goals

- read simple inductive, structure, and class declarations;
- predict the cases of a pattern match from the constructors of a type;
- recognize the shape of recursive definitions and induction proofs;
- use basic examples of `Option`, `List`, `Fin`, `Multiset`, and `Finset`;
- understand what instance search and deriving are for;
- explain the roles of `Repr`, `BEq`, `DecidableEq`, and `Inhabited`;
- distinguish explicit arguments, implicit arguments, and type class arguments;
- read mathematical class assumptions such as `[Group G]`, `[Ring R]`, and `[Field K]`.

Four-Part Outline

Part I	Inductive Motivation and basic inductive declarations
Part II	Recursion and Induction Pattern matching, recursion, cases, and induction
Part III	Structures and Models structure, Option, Inhabited, List, Fin, Multiset, Finset
Part IV	Type Classes class, instance, #synth, Group, Ring, deriving, and Expr

Why Inductive Types?

Many Lean objects are built from constructors:

- Bool: true, false
- Nat: zero, succ
- Option alpha: none, some a
- List alpha: [], x :: xs
- Knowing how values are built tells us how to compute with them and how to prove things about them.

A First Example

```
inductive Weekday where
  | mon | tue | wed | thu | fri | sat | sun
deriving Repr, BEq, DecidableEq
```

- Weekday is a new type.
- The constructors are the only ways to build values of this type.
- deriving asks Lean to generate useful standard behavior.

How to Read an Inductive Declaration

Ask four questions:

- ① What is the new type?
- ② What are the constructors?
- ③ Which constructors carry data?
- ④ Which constructors are recursive?

Constructor shape predicts program shape and proof shape.

Constructors Define All Cases

```
inductive Weekday where  
  | mon | tue | wed | thu | fri | sat | sun
```

- Every Weekday is one of these seven constructors.
- A function out of Weekday should account for these seven cases.
- A proof about all weekdays can also proceed by these seven cases.

Live Lean: `codeD2/D2_Inductive.lean`

Constructors Are Functions

```
inductive MyNat where  
  | zero : MyNat  
  | succ : MyNat -> MyNat
```

- zero is already a natural number.
- succ needs a smaller natural number before it produces a new one.
- Use #check to inspect constructor types.

Parameters vs. Constructor Arguments

```
inductive MyOption (alpha : Type) where  
  | none : MyOption alpha  
  | some : alpha -> MyOption alpha
```

- alpha is a parameter of the whole family of types.
- some stores one particular value of type alpha.
- Examples: MyOption Nat, MyOption String.

Recursive Constructors

```
inductive MyNat where
  | zero : MyNat
  | succ : MyNat -> MyNat

inductive MyList (alpha : Type) where
  | nil : MyList alpha
  | cons : alpha -> MyList alpha -> MyList alpha
```

- Recursive constructors mention the type being defined.
- The recursive argument is the smaller object.
- This is why structural recursion and induction follow the same shape.

Recursive Constructors: Natural Numbers

```
def MyNat.add : MyNat -> MyNat -> MyNat
  | m, zero => m
  | m, succ n => succ (MyNat.add m n)
```

- The definition recurses on the second argument.
- The two equations follow the two constructors of `MyNat`.
- This is the same pattern as recursive definitions in mathematics.

Trees: Two Recursive Subobjects

```
inductive Tree (alpha : Type) where
  | leaf : alpha -> Tree alpha
  | node : Tree alpha -> Tree alpha -> Tree alpha
```

- A leaf stores one value.
- A node stores two subtrees.
- Later, the node case of an induction proof will have two induction hypotheses.

Data shape $\longrightarrow$ program shape $\longrightarrow$ proof shape

We now use constructors in two ways:

- pattern matching to inspect data;
- recursion to compute over recursive data.

Pattern Matching

```
def isWeekend : Weekday -> Bool
| Weekday.sat => true
| Weekday.sun => true
| _ => false
```

- Pattern matching eliminates an inductive value.
- Each branch handles one possible constructor shape.
- The wildcard pattern `_` handles remaining cases.

Live Lean: `codeD2/D2_PatternMatching.lean`

Matching on Option

```
def getWithDefault {alpha : Type}  
  (default : alpha) : Option alpha -> alpha  
  | none => default  
  | some a => a
```

- The none branch has no stored value.
- The some branch exposes the stored value.
- The return type is the same in both branches.

Recursive Definitions on Natural Numbers

```
def add (m : Nat) : Nat -> Nat
| 0 => m
| n + 1 => add m n + 1

def mul (m : Nat) : Nat -> Nat
| 0 => 0
| n + 1 => mul m n + m
```

- The recursion is on the second argument.
- The equations are ordinary recursive mathematical definitions.
- Powers can be defined in the same style.

Recursive Definitions on Lists

```
def myLength {alpha : Type} : List alpha -> Nat
| [] => 0
| _ :: xs => myLength xs + 1

def myMap {alpha beta : Type} (f : alpha -> beta) :
  List alpha -> List beta
| [] => []
| x :: xs => f x :: myMap f xs
```

- Both definitions branch on `[]` and `x :: xs`.
- The recursive call is made on `xs`, a smaller list.
- `map` preserves the list shape while changing stored data.

From Recursion to Induction

Data	Program	Proof
base constructor	base branch	base case
recursive constructor	recursive call	induction hypothesis

Programs and proofs follow the same shape as the data.

Proof by Cases

```
example (b : Bool) : b = true \/\ b = false := by
  cases b <;> simp
```

- cases splits a value according to its constructors.
- It is best for small or non-recursive inductive types.

Live Lean: `codeD2/D2_InductionProofs.lean`

Induction on Lists

```
theorem append_nil_right {alpha : Type} (xs : List alpha) :  
  xs ++ [] = xs := by  
  induction xs with  
  | nil =>  
    rfl  
  | cons x xs ih =>  
    rw [ih]
```

- The base case is the empty list.
- The step case is $x :: xs$.
- The induction hypothesis is the theorem already known for xs .

Induction on Addition

```
theorem zero_add (n : Nat) : add 0 n = n := by
  induction n with
  | zero =>
    rfl
  | succ n ih =>
    calc
      add 0 (n + 1) = add 0 n + 1 := rfl
      _ = n + 1 := by rw [ih]
```

- The base case proves the statement for 0.
- The step case uses the induction hypothesis for n .
- This is the familiar proof of a recursive equation by induction.

Reading an Induction Proof

In an induction step, focus on:

- the variables introduced by the constructor;
- the smaller recursive object;
- the induction hypothesis;
- the goal after simplification.

The local context is not decoration; it tells us exactly what we are allowed to use.

20-minute break

We now move from raw data shapes to reusable structure:

- `structure`: bundle named fields;
- common inductive types: choose the right model;
- `class`: bundle structure that Lean can find automatically.

Structures as Records

```
structure Vector2 where  
  x : Int  
  y : Int
```

- A structure bundles named fields.
- Values can be constructed with record syntax.
- Fields can be accessed with dot notation.

Live Lean: `codeD2/D2_Structure.lean`

Operations on a Structure

```
def Vector2.add (u v : Vector2) : Vector2 :=  
  { x := u.x + v.x, y := u.y + v.y }  
  
def Vector2.dot (u v : Vector2) : Int :=  
  u.x * v.x + u.y * v.y
```

- `u.x` and `u.y` access coordinates.
- Record syntax constructs new vectors from field values.
- This mirrors ordinary coordinate definitions in linear algebra.

Bundling Data with a Property

```
structure PointOnDiagonal where  
  x : Int  
  y : Int  
  onDiagonal : x = y
```

- Structures can contain data fields and proof fields.
- The field `onDiagonal` records a mathematical property.
- This pattern is common in formalized mathematics.

Inductive vs. Structure

Tool	Typical use	Shape
inductive	alternatives, recursive data	one or more constructors
structure	bundled fields	one constructor with named fields

Many mathematical objects in Lean are represented as structures.

Common Types as Modeling Choices

Type	Informal meaning	Typical use
<code>Option alpha</code>	maybe a value	partial lookup
<code>List alpha</code>	ordered finite data	sequences, recursion
<code>Fin n</code>	a natural number below n	finite indices
<code>Multiset alpha</code>	unordered, with multiplicity	combinatorics
<code>Finset alpha</code>	unordered, no duplicates	finite sets and sums

Live Lean: `codeD2/D2_CommonTypes.lean`

Option: Handling Failure

```
def safeHead {alpha : Type} : List alpha -> Option alpha
| [] => none
| x :: _ => some x
```

- The type tells us that lookup may fail.
- Any consumer must handle both `none` and `some`.

Option: Supplying a Default

```
def getWithDefault {alpha : Type}
  (default : alpha) : Option alpha -> alpha
| none => default
| some x => x

#eval getWithDefault 0 (some 37)
#eval getWithDefault 0 (none : Option Nat)
```

- The caller explicitly provides the fallback value.
- The type `alpha` may be arbitrary.
- Without extra information, Lean cannot invent an `alpha`.

Option with an Inhabited Default

```
#check Inhabited
#check default
#synth Inhabited Nat

def getOrElseDefault {alpha : Type} [Inhabited alpha] :
  Option alpha -> alpha
  | none => default
  | some x => x

#eval getOrElseDefault (none : Option Nat)
#eval getOrElseDefault (none : Option Bool)
#eval getOrElseDefault (none : Option (List Nat))
```

- `[Inhabited alpha]` asks Lean to find a default.
- The square brackets mark a type class parameter.
- This is the same search mechanism used later for algebraic structure.

What Inhabited Does Not Mean

- `[Inhabited alpha]` means: Lean, please find an inhabitant of `alpha`.
- It does not mean that the default is mathematically unique.
- It does not mean that the default is the best modeling choice.
- It only means that a value can be found automatically when requested.

List: Ordered Finite Data

For mathematicians, `List alpha` is often best read as a finite sequence.

Common operations:

- `map`: transform each element;
- `filter`: keep elements satisfying a test;
- `foldl`, `foldr`: accumulate values;
- `length`, `append`, `membership`.

Order matters, and repetitions are allowed.

Fin: Safe Bounded Indexing

```
#check Fin 3  
#check (0 : Fin 3)  
#check (2 : Fin 3)
```

- An element of `Fin n` is a natural number together with a proof it is less than n .
- Read it as the formal index set $\{0, \dots, n - 1\}$.
- Invalid indices are rejected by the type system.
- This is useful for vectors, matrices, and finite choices.

Finite Data: Three Choices

Type	Order	Repetition
List α	kept	kept
Multiset α	ignored	kept
Finset α	ignored	removed

- Use `List` for finite sequences and recursive traversals.
- Use `Multiset` when multiplicity matters but order does not.
- Use `Finset` for finite sets, finite sums, and finite products.

Multiset and Finset in Mathlib

```
import Mathlib

def orderedData : List Nat := [1, 2, 1]
def unorderedWithMultiplicity : Multiset Nat := [1, 2, 1]
def finiteSet : Finset Nat := {1, 2, 1}

#eval orderedData.length
#eval unorderedWithMultiplicity.card
#eval unorderedWithMultiplicity.count 1
#eval finiteSet.card
```

Live Lean: `codeD2/D2_Typeclass_Mathlib.lean`

From Structures to Classes

- A structure bundles data and operations.
- A class is a structure whose instances Lean can search for automatically.
- In mathematics, this is how notation and algebraic structure are reused across many types.

A type class parameter says: Lean, please find the required structure.

Why Type Classes?

Type classes let us reuse definitions and theorems across many types.

- `Nat`, `Int`, `Rat`, and `Complex` all support addition.
- Many structures share notation such as `+`, `*`, `0`, and `1`.
- A theorem should often ask for the structure it needs, not for one fixed type.

Do not ask: is this exactly `Nat` or `Int`? Ask: does this type have the structure I need?

Class as Searchable Structure

A class is very close to a structure:

- it has fields;
- values of the class are pieces of structure;
- instances register those pieces of structure for later search.

The key extra behavior is not mathematical magic. It is automatic lookup.

A Small Type Class

```
class HasNorm (alpha : Type) where
  norm : alpha -> Nat

def doubleNorm [HasNorm alpha] (x : alpha) : Nat :=
  2 * HasNorm.norm x
```

- `[HasNorm alpha]` is a type class parameter.
- Lean tries to synthesize the instance automatically.

Live Lean: `codeD2/D2_Typeclass.lean`

What Is the Instance?

```
class HasNorm (alpha : Type) where  
  norm : alpha -> Nat
```

An instance for `HasNorm alpha` is a value containing:

- a function `norm : alpha -> Nat`;
- registered so that Lean can find it later.

The class says what data is required. The instance supplies that data.

```
structure Vector2 where
  x : Int
  y : Int

instance : HasNorm Vector2 where
  norm v := intAbsNat v.x + intAbsNat v.y
```

- An instance provides the structure required by a class.
- Once registered, it can be found by instance search.
- #synth lets us inspect what Lean can find.

Checking Instance Search

```
#synth HasNorm Vector2
#synth HasNorm Nat

def natNormInstance : HasNorm Nat :=
  inferInstance

#eval natNormInstance.norm 9
```

- #synth checks instance search at command level.
- inferInstance performs instance search inside a term.
- Both are useful for understanding what Lean can find.

Explicit Structure Argument

```
def useExplicit {alpha : Type}
  (inst : HasNorm alpha) (x : alpha) : Nat :=
  inst.norm x

#eval useExplicit (inferInstance : HasNorm Nat) 5
```

- The instance is an ordinary explicit argument.
- The caller must supply it.
- This is honest, but quickly becomes verbose.

Type Class Argument

```
def useInstance {alpha : Type}
  [HasNorm alpha] (x : alpha) : Nat :=
  HasNorm.norm x

#eval useInstance (5 : Nat)
```

- The square brackets activate instance search.
- Lean infers $\alpha = \text{Nat}$ from the argument.
- Then it searches for HasNorm Nat .

Three Kinds of Parameters

Syntax	Meaning	Who supplies it?
<code>(x : alpha)</code>	explicit value	user writes it
<code>{alpha : Type}</code>	ordinary implicit parameter	elaborator infers it
<code>[HasNorm alpha]</code>	type class parameter	instance search finds it

Type class parameters are still parameters. They are just filled by a special search procedure.

Named Instance Parameters

```
def useInstanceVerbose {alpha : Type}
  [inst : HasNorm alpha] (x : alpha) : Nat :=
  inst.norm x
```

- We can name the instance if we want to use it directly.
- The brackets still request instance search.
- This is useful when several fields or projections are needed.

When Search Fails

```
class HasSize (alpha : Type) where
  size : alpha -> Nat

instance : HasSize String where
  size s := s.length

#synth HasSize String
-- #synth HasSize Nat
```

- The class HasSize exists.
- Lean can find an instance for String.
- Lean cannot find an instance for Nat unless one is registered.

Notation Comes from Classes

```
#synth Add Nat
#synth Add Int
#synth Mul Nat
#synth Mul Int

def addThree [Add alpha] (a b c : alpha) : alpha :=
  a + b + c
```

- The symbol `+` is interpreted using an `Add` instance.
- The symbol `*` is interpreted using a `Mul` instance.
- The same notation can work over many mathematical types.

How Lean Reads Addition

When Lean sees:

$$a + b$$

it must know:

- the type of a and b ;
- an `Add` instance for that type;
- how the notation expands using that instance.

The same surface notation can mean different operations at different types.

Generic Notation

```
def addThree {alpha : Type} [Add alpha]
  (a b c : alpha) : alpha :=
  a + b + c

#eval addThree (1 : Nat) 2 3
#eval addThree (1 : Int) 2 3
```

- The definition is generic in `alpha`.
- The operation comes from `[Add alpha]`.
- The same code runs at different concrete types.

Mathlib: Algebraic Instances

```
import Mathlib

#synth Add Nat
#synth Add Int
#synth Add Rat
#synth Add Complex

#synth AddCommMonoid Nat
#synth AddCommGroup Int
#synth Ring Int
#synth Field Rat
#synth Field Complex
```

- Mathlib provides a rich algebraic hierarchy.
- Lean can find instances for familiar number systems.
- This is why generic algebraic theorems can apply broadly.

Live Lean: `codeD2/D2_Typeclass_Mathlib.lean`

Classes for Mathematical Structure

Class	Purpose
Add, Mul	notation for operations
Zero, One	distinguished elements
AddCommMonoid	addition with 0, associativity, commutativity
Group, AddGroup	inverses, multiplicative or additive notation
Ring, CommRing	ring structure
Field	field structure

Structure, Class, Instance

Word	Reading
structure	bundle a concrete object with named fields
class	bundle reusable structure on a type
instance	register that a type carries this structure
[Group G]	ask Lean to find a group structure on G

This is why theorems can be stated once and reused for many mathematical objects.

Multiplicative and Additive Notation

Class	Identity	Operation	Inverse
[Group G]	1	$g * h$	g^{-1}
[AddGroup A]	0	$a + b$	$-a$

The mathematical idea is similar, but the notation and class names are different.

Generic Group Theorems

```
example {G : Type} [Group G] (g h : G) :  
  g * h * h^-1 = g := by  
  simp [mul_assoc]  
  
example {A : Type} [AddGroup A] (a b : A) :  
  a + b + (-b) = a := by  
  simp [add_assoc]
```

- The type is arbitrary.
- The class assumption supplies notation and algebraic facts.
- `simp` can use lemmas from the instance.

Operations vs. Laws

- `[Add alpha]` gives an addition operation.
- It does not by itself state associativity, commutativity, or identity laws.
- For laws, use stronger classes such as `AddCommMonoid alpha`, `Group G`, `Ring R`, or `Field K`.

A theorem with `[Ring R]` applies to any type for which Lean can find a ring instance.

Data Classes and Property Classes

Kind	Example	What the instance contains
Data class	<code>HasNorm alpha</code>	a function <code>norm : alpha -> Nat</code>
Data class	<code>Add alpha</code>	an operation <code>add : alpha -> alpha -> alpha</code>
Property class	<code>MulCommLaw alpha</code>	a proof of commutativity
Mixed structure	<code>Group G</code>	operations plus laws

In mathlib, algebraic classes usually contain both data fields and proof fields.

A Prop-Valued Class

```
class MulCommLaw (alpha : Type) [Mul alpha] : Prop where
  mul_comm_law : forall a b : alpha, a * b = b * a

instance : MulCommLaw Nat where
  mul_comm_law := Nat.mul_comm

example {alpha : Type} [Mul alpha] [MulCommLaw alpha]
  (a b : alpha) : a * b = b * a :=
  MulCommLaw.mul_comm_law a b
```

Live Lean: `codeD2/D2_Typeclass.lean`

When a Math Structure Does Not Exist

```
-- Nat has addition, but no additive inverses:  
-- #synth AddGroup Nat  
  
-- Int is a ring, but not a field:  
-- #synth Field Int
```

- These are useful failures for teaching instance search.
- The class names exist, but Lean cannot find the requested instances.
- This is different from an unknown identifier error.

Deriving Standard Instances

```
inductive Color where  
  | red | green | blue  
deriving Repr, BEq, DecidableEq
```

- deriving generates standard instances.
- It makes custom types usable in small examples immediately.

Live Lean: `codeD2/D2_Deriving.lean`

Class	Purpose	Typical use
Repr alpha	printable representation	<code>#eval</code> output
BEq alpha	Boolean equality test	<code>x == y</code>
DecidableEq alpha	decidable equality proposition	<code>if x = y then ...</code>

These are ordinary type class instances, often generated automatically.

Repr: Printing Values

```
inductive Color where
  | red | green | blue
  deriving Repr

#eval Color.red
```

- Repr lets Lean display values.
- It is especially useful for live demos.
- It is not a mathematical theorem; it is display infrastructure.

BEq: Boolean Equality

```
inductive Direction where
  | north | south | east | west
  deriving Repr, BEq

#eval Direction.north == Direction.north
#eval Direction.north == Direction.south
```

- BEq provides ==.
- The result is a Bool.
- This is computational equality testing.

DecidableEq: Equality as a Proposition

```
inductive Suit where
  | clubs | diamonds | hearts | spades
  deriving Repr, BEq, DecidableEq

example : DecidableEq Suit := by
  infer_instance
```

- DecidableEq decides propositions $x = y$.
- This is different from returning a Boolean.
- It supports proof-oriented uses of equality.

BEq vs DecidableEq

Expression	Type	Meaning
<code>x == y</code>	<code>Bool</code>	computes a Boolean answer
<code>x = y</code>	<code>Prop</code>	states a proposition
<code>Decidable (x = y)</code>	<code>Type</code>	gives a decision procedure

In practice, both are useful, but they serve different layers of Lean.

What Deriving Does

- It inspects the constructors of an inductive type or the fields of a structure.
- It generates instance declarations.
- Those instances are then available to `#eval`, `==`, `decide`, and library APIs.

Connection to Part I

The constructor shape is enough for Lean to generate many standard behaviors.


The extends Keyword

`extends` says that a class includes the fields and instance behavior of parent classes.

- It is how Lean builds class hierarchies.
- A child class can be used wherever a parent class is required.
- This is essential for algebraic hierarchies.
- It also creates design risks when parent data is duplicated.

A Diamond Warning

Instance search is powerful, but class hierarchies must be designed carefully.



A value of `Both` may carry structure through two paths.

What Is the Diamond Problem?

- Two parent classes may contain or inherit what is mathematically intended to be the same structure.
- Lean sees data, not intention.
- If two copies are stored, Lean will not silently identify them.
- Mathlib's hierarchy is carefully designed to avoid accidental duplication.

Practical advice

Prefer existing mathlib hierarchies over inventing parallel class towers.

A Mathematical Diamond Example

```
class BadSemigroup (alpha : Type) where
  op : alpha -> alpha -> alpha
class BadAddSemigroup (alpha : Type)
  extends BadSemigroup alpha
class BadMulSemigroup (alpha : Type)
  extends BadSemigroup alpha
instance : BadAddSemigroup Int where
  op := Int.add
instance : BadMulSemigroup Int where
  op := Int.mul
#check BadAddSemigroup.toBadSemigroup
#check BadMulSemigroup.toBadSemigroup
#eval (BadAddSemigroup.toBadSemigroup (alpha := Int)).op 2 3
#eval (BadMulSemigroup.toBadSemigroup (alpha := Int)).op 2 3

class BadRing (alpha : Type)
  extends BadAddSemigroup alpha, BadMulSemigroup alpha
-- Failing attempt:
-- instance : BadRing Int where
--   toBadAddSemigroup := inferInstance
--   toBadMulSemigroup := inferInstance
```

- $(\mathbb{Z}, +)$ and $(\mathbb{Z}, *)$ are both semigroup-like.
- `extends` creates projections to the parent class.
- The failing instance tries to merge two different parent `op` fields.

Fixing the Diamond

```
class GoodAddSemigroup (alpha : Type) where
  add : alpha -> alpha -> alpha

class GoodMulSemigroup (alpha : Type) where
  mul : alpha -> alpha -> alpha

class GoodRingSkeleton (alpha : Type)
  extends GoodAddSemigroup alpha, GoodMulSemigroup alpha

instance : GoodRingSkeleton Int where
  add := Int.add
  mul := Int.mul

#eval (inferInstance : GoodRingSkeleton Int).add 2 3
#eval (inferInstance : GoodRingSkeleton Int).mul 2 3
```

- Use separate fields for mathematically different operations.
- Mathlib follows this idea: AddSemigroup uses $+$, Semigroup uses $*$.

Diamond: Practical Reading

- For ordinary use, trust mathlib's existing hierarchies.
- Avoid inventing parallel algebraic class towers in class projects.
- If two paths should share data, design the shared data as a common parent.
- If two fields should be equal, add that equality as a law.

The important lesson is not fear of type classes, but respect for hierarchy design.

Expression Trees

```
inductive Expr where
  | const : Nat -> Expr
  | var : String -> Expr
  | add : Expr -> Expr -> Expr
  deriving Repr, BEq, DecidableEq
```

Syntax trees are naturally inductive.

Live Lean: `codeD2/D2_IntegratedExample.lean`

Why Expr Is a Good Final Example

Expr combines almost everything from today:

- constructors define a syntax tree;
- recursive constructors create recursive programs;
- `Option` handles failed lookup;
- `List` stores collected variables;
- structure packages environments;
- deriving gives printing and equality;
- induction proves properties of tree transformations.

Today this is only a first impression of expression trees.

An Environment

```
structure Env where
  lookup : String -> Option Nat

def sampleEnv : Env :=
  { lookup := fun name =>
    if name = "x" then some 10
    else if name = "y" then some 5
    else none }
```

- An environment interprets variable names.
- Lookup can fail, so the return type is `Option Nat`.
- This is a structure bundling one operation.

Evaluating Expressions

```
def Expr.eval (env : Env) : Expr -> Option Nat
| const n => some n
| var name => env.lookup name
| add left right =>
    match eval env left, eval env right with
    | some m, some n => some (m + n)
    | _, _ => none
```

- Evaluation is recursive on the expression tree.
- Addition succeeds only if both subexpressions succeed.
- Failure propagates through none.

Collecting Variables

```
def Expr.vars : Expr -> List String
  | const _ => []
  | var name => [name]
  | add left right => vars left ++ vars right
```

- This is another recursive traversal of the same tree.
- List stores finite ordered data.
- The constructor shape again determines the program shape.

Mirroring an Expression Tree

```
def Expr.mirror : Expr -> Expr
  | const n => const n
  | var name => var name
  | add left right => add (mirror right) (mirror left)
```

- Constants and variables stay fixed.
- Binary nodes swap their two subtrees.
- The recursive calls go to smaller subexpressions.

A Proof About Mirror

```
theorem mirror_mirror (e : Expr) :  
  Expr.mirror (Expr.mirror e) = e := by  
  induction e with  
  | const n => rfl  
  | var name => rfl  
  | add left right ihLeft ihRight =>  
    simp [Expr.mirror, ihLeft, ihRight]
```

Reading the Expr Induction Proof

- There is one case for each constructor.
- Constant and variable cases are immediate.
- Addition has two induction hypotheses.

This is the tree version of the same pattern we saw for lists.

A Glimpse of Lean's Real Expr

```
inductive Lean.Expr where
| bvar      : Nat -> Expr
| fvar      : FVarId -> Expr
| mvar      : MVarId -> Expr
| sort      : Level -> Expr
| const     : Name -> List Level -> Expr
| app       : Expr -> Expr -> Expr
| lam       : Name -> Expr -> Expr -> BinderInfo -> Expr
| forallE   : Name -> Expr -> Expr -> BinderInfo -> Expr
| letE      : Name -> Expr -> Expr -> Expr -> Bool -> Expr
| lit       : Literal -> Expr
| mdata     : MData -> Expr -> Expr
| proj      : Name -> Nat -> Expr -> Expr
```

Why Lean.Expr Matters

```
import Lean

#check Lean.Expr
#check Lean.Expr.bvar
#check Lean.Expr.const
#check Lean.Expr.app
#check Lean.Expr.lam
#check Lean.Expr.forallE
#check Lean.Expr.letE
```

- Our Expr is a small teaching language.
- Lean's real Expr represents terms in Lean itself.
- We only need the impression, not elaboration or metaprogramming.

- Today: expressions as inductive syntax trees.
- Later: functional programs over such trees.
- Later: transformations, optimizations, and correctness proofs.
- The same recursion and induction principles will keep appearing.

Expr as Integration Point

With Expr, we can use:

- inductive to define syntax trees;
- pattern matching to evaluate or analyze expressions;
- Option for failed variable lookup;
- List to collect variable names;
- structure to bundle an environment;
- deriving for printing and equality;
- induction to prove properties of recursive functions.

Suggested Exercises

- Complete the guided holes in `D2_Exercises.lean`.
- Use `Inhabited` to write a default-returning `Option` function.
- Use `inferInstance` to fill an explicit class argument.
- Experiment with `Repr`, `BEq`, and `DecidableEq`.
- Modify the simple `Expr` language.
- Extend the evaluator and variable collector.
- Prove a small theorem by induction on expressions.

Live Lean: `codeD2/D2_Exercises.lean`